

European rainfall quantile mapping exercise

Jordan Richards

2023-11-17

Dataset

- ▶ The data consists of:
 - ▶ hourly summer rainfall (mm) over 500 grid-boxes (0.25deg by 0.25deg) across Western Europe (2021-2022).
 - ▶ we will treat grid-box as point location.
- ▶ collection of 11 relevant meteorological and orographical covariates at each site.
- ▶ Obtained from ERA5:
 - ▶ <https://cds.climate.copernicus.eu/cdsapp#!/dataset/reanalysis-era5-single-levels>
 - ▶ we will use a GNN model to map high quantiles of hourly rainfall, as well as estimate the probability of extreme rainfall at a single site.

Load data

```
load("../Data/EuroRain.Rdata")  
dim(Y)
```

```
## [1] 4416 500
```

```
dim(X)
```

```
## [1] 4416 500 11
```

Predictors

- ▶ The predictors are:
 - ▶ meteorological: air temperature at 2m, mean sea level pressure, surface pressure, ozone, 10m u- and v-wind speed components
 - ▶ orographical: angle, isotropy, land-sea mask, slope, standard deviation

```
print(cov_names[1:5])
```

```
## [1] "t2m" "msl" "sp" "tco3" "u10"
```

```
print(cov_names[5:11])
```

```
## [1] "u10" "v10" "anor" "isor" "lsm" "slor" "sdor"
```

Data normalisation

- ▶ We first scale the input data to improve the numerical stability of training

```
#Normalise inputs
X_scaled <- X
for(i in 1:dim(X)[3]){
  temp <- X[, ,i]
  m <- mean( temp, na.rm=T)
  s <- sd( temp, na.rm=T)
  temp <- ( temp-m)/s
  X_scaled[, ,i] <- temp
}
```

Get some validation data

```
#Normalise inputs
```

```
set.seed(1)
```

```
validation.inds <- sample(1:nrow(Y), nrow(Y)/5) #20%
```

```
X.valid <- X_scaled[validation.inds,,]
```

```
Y.valid <- Y[validation.inds,]
```

```
X.train <- X_scaled[-validation.inds,,]
```

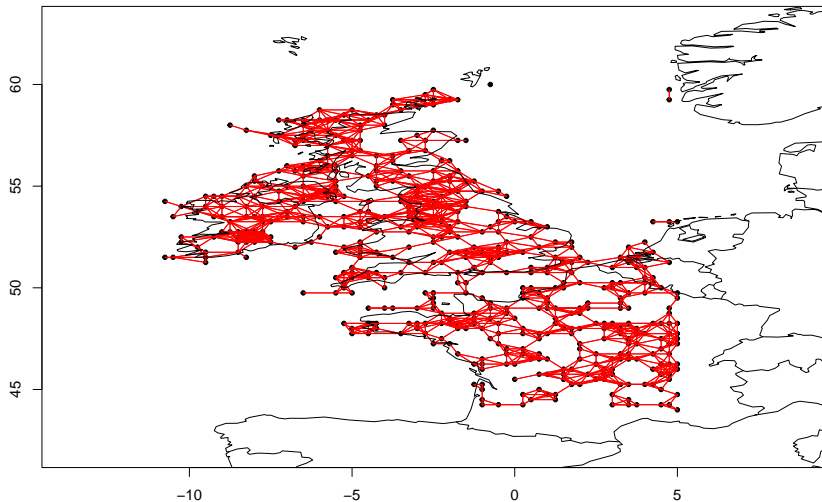
```
Y.train <- Y[-validation.inds,]
```


Building a graph

- ▶ We can build a graph for our data using the location coordinates. The graph structure will be stored in an adjacency matrix. This matrix will have entries 1 if sites are within 75km of each other.

```
dist<-fields::rdist.earth(coords,miles=F)
cut.off.dist <- 75
A <- dist
A[dist>cut.off.dist]=0; A[dist <= cut.off.dist]=1
diag(A)=0
```


Building a graph



Graph layers

```
spk <- reticulate::import("spektral")  
print(names(sp$layers)[1:9])
```

```
[1] "activations" "agnn_conv" "AGNNConv" "appnp_conv"  
"APPNPConv"
```

```
[6] "arma_conv" "ARMAConv" "base" "censnet_conv"
```

We can create a custom Keras layer with an embedded spektral layer. Here we used a graph convolutional layer with trainable skip connection.

See also <https://graphneural.network/layers/convolution/>

Custom graph layers

```
layer_graph_conv <- function(  
  object,  
  channels,  
  activation = NULL,  
  use_bias = TRUE,  
  kernel_initializer = 'glorot_uniform',  
  bias_initializer = 'zeros',  
  kernel_regularizer = NULL,  
  bias_regularizer = NULL,  
  activity_regularizer = NULL,  
  kernel_constraint = NULL,  
  bias_constraint = NULL,  
  name=NULL,  
  ...)  
{... #REPLACE WITH FOLLOWING}
```

Custom graph layers (2)

```
args <- list(channels = as.integer(channels),
             activation = activation,
             use_bias = use_bias,
             kernel_initializer = kernel_initializer,
             bias_initializer = bias_initializer,
             kernel_regularizer = kernel_regularizer,
             bias_regularizer = bias_regularizer,
             activity_regularizer = activity_regularizer,
             kernel_constraint = kernel_constraint,
             bias_constraint = bias_constraint,
             name=name
)
keras::create_layer(spk$layers$GCSCnv, object, args)
```

Keras GNN

We can now build a Keras model as we have done previously, just by swapping `layer_dense` (or equivalent) with `layer_graph_conv`.

However, we need an extra input to the graph layer that includes information about the graph itself. This extra input depends on the type of layer. For `GCGSconv`, we need the normalised adjacency matrix.

```
ML<-spk$utils$convolution$normalized_adjacency(A)
```

Keras GNN

```
library(keras)

input.lay <- layer_input(shape = dim(X)[2:3])

hidden.lay1<- list(input.lay, ML) %>%
  layer_graph_conv(channels = 32, activation = "relu")

hidden.lay2<- list(hidden.lay1, ML) %>%
  layer_graph_conv(channels = 32, activation = "relu")
```

Keras GNN

```
library(keras)
# I want strictly positive quantiles only,
# so lets apply an exponential transformation
output.lay<- hidden.lay2 %>%
  layer_dense(units = 1,   activation = "exponential")

model <- keras_model(
  inputs = c(input.lay),
  outputs = c(output.lay)
)
```

Compiling

We will compile with the adam optimiser and our quantile loss function

```
model %>% compile(  
  loss = tilted_loss,  
  optimizer = "adam"  
)
```


Fitting the model

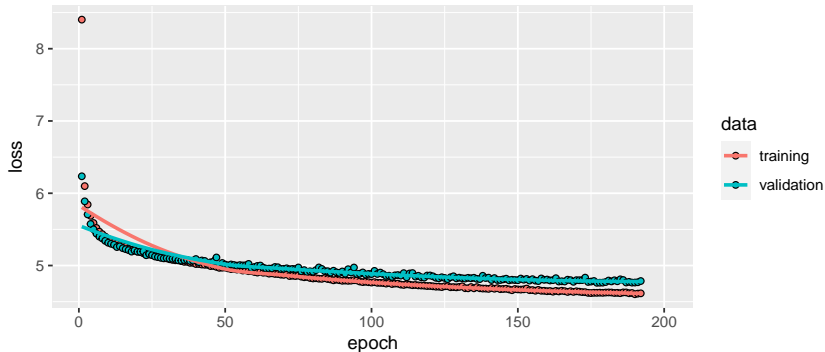
I'll also shuffle the training data to break up any temporal dependence. This will help the training procedure.

```
history <- model %>% fit(  
  x = X.train, y = Y.train, shuffle=T,  
  epochs = 200, batch_size = 32,  
  validation_data = list(  
    X.valid, Y.valid),  
  callbacks = list(  
    callback_model_checkpoint(filepath = "model_weights",  
      verbose=0, monitor="val_loss",  
      save_best_only = TRUE, save_weights_only = TRUE),  
    callback_early_stopping(monitor = "val_loss",  
      min_delta = 0, patience = 15)),  
  verbose=0  
)  
model <- load_model_weights_tf(model,  
  filepath="model_weights")
```

Fitting the model

- Check for overfitting

```
plot(history)
```



Predictions

```
qhat <- model %>% predict(X_scaled)
```

```
## 138/138 - 1s - 1s/epoch - 10ms/step
```

```
mean(qhat[,1] > Y)
```

```
## [1] 0.8925417
```

Weighted adjacency

- ▶ Let's try weighting the adjacency matrix, and see if that provides better predictions. We will use a Gaussian weighting kernel with range 50

```
A2 <- exp(-(dist/50)^2)
diag(A2) <- 0
A2[dist > cut.off.dist] <- 0
ML2<-spk$utils$convolution$normalized_adjacency(A2)
```

New model

```
library(keras)
input.lay2 <- layer_input(shape = dim(X)[2:3])
hidden.lay21<- list(input.lay2, ML2) %>%
  layer_graph_conv(channels = 32, activation = "relu")
hidden.lay22<- list(hidden.lay21, ML2) %>%
  layer_graph_conv(channels = 32, activation = "relu")
output.lay2 <- hidden.lay22 %>%
  layer_dense(units=1, activation = 'exponential')

model2 <- keras_model(
  inputs = c(input.lay2),
  outputs = c(output.lay2)
)
```

New model fitting

```
model2 %>% compile(  
  loss = tilted_loss,  
  optimizer = "adam"  
)
```

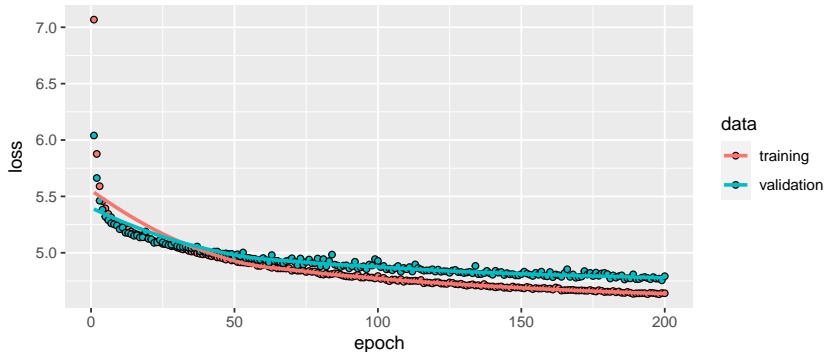
New model fitting

```
history2 <- model2 %>% fit(  
  x = X.train, y = Y.train, shuffle=T,  
  epochs = 200, batch_size = 32,  
  validation_data = list(  
    X.valid, Y.valid),  
  callbacks = list(  
    callback_model_checkpoint(filepath = "model_weights",  
      verbose=0, monitor="val_loss",  
      save_best_only = TRUE, save_weights_only = TRUE),  
    callback_early_stopping(monitor = "val_loss",  
      min_delta = 0, patience = 15)),  
  verbose=0  
)  
model2 <- load_model_weights_tf(model2,  
  filepath="model_weights")
```

Fitting the model

- Check for overfitting

```
plot(history2)
```



Comparison

```
qhat2 <- model2 %>% predict(X_scaled)
```

```
## 138/138 - 2s - 2s/epoch - 11ms/step
```

```
mean(qhat2[, , 1] >= Y)
```

```
## [1] 0.8959579
```

```
print(tilted_loss(Y.valid, predict(model, X.valid)[ , , 1]))
```

```
## 28/28 - 0s - 493ms/epoch - 18ms/step
```

```
## tf.Tensor(4.760883725181087, shape=(), dtype=float64)
```

```
print(tilted_loss(Y.valid, predict(model2, X.valid)[ , , 1]))
```

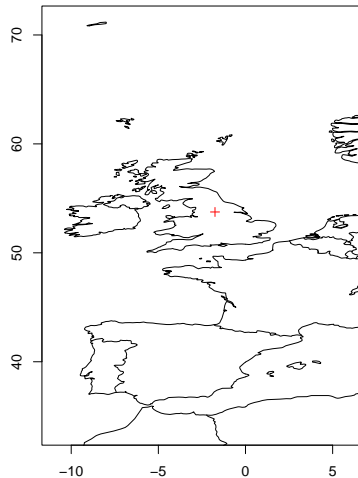
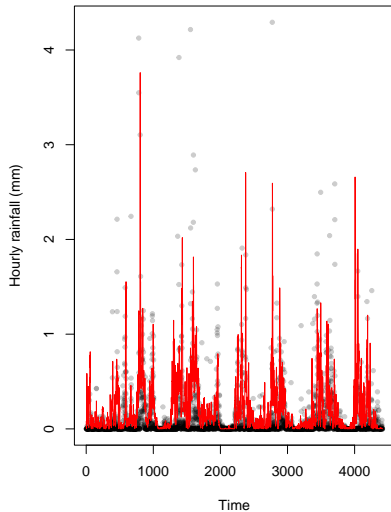
```
## 28/28 - 0s - 485ms/epoch - 17ms/step
```

```
## tf.Tensor(4.755669896862342, shape=(), dtype=float64)
```

Lower validation loss when using the weighted A matrix.

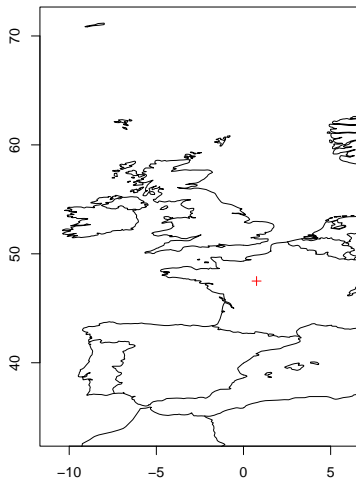
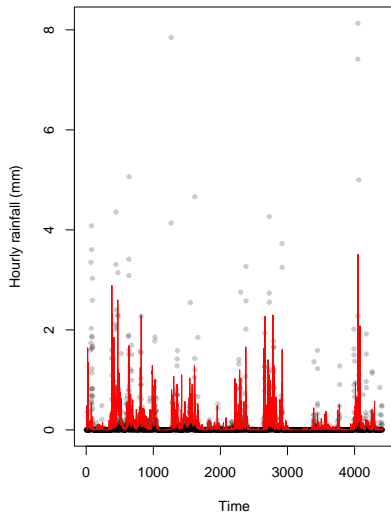
Predictions

► Check predictions



Predictions

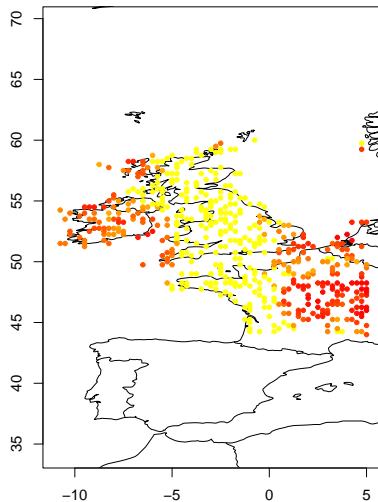
► Check predictions



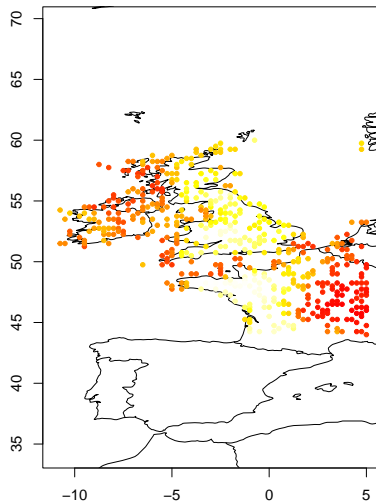
Prediction maps

Hourly rainfall (mm)

Quantile (mm)



2021-06-05 06:00:00

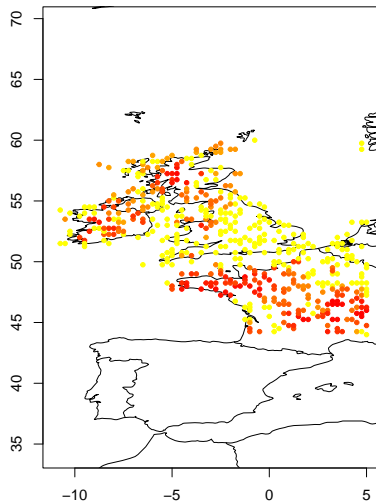


2021-06-05 06:00:00

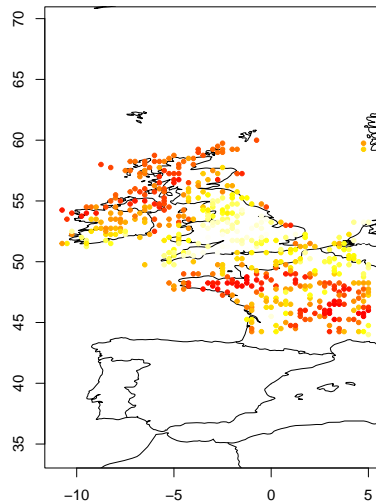
Prediction maps

Hourly rainfall (mm)

Quantile (mm)



2022-08-14 18:00:00



2022-08-14 18:00:00